

+ Code + Text

#21162101007\_KSHITIJGUPTA

PR5

```
[ ] import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, precision_recall_fscore_support, log_loss, classification_report, accuracy_score

[ ] from google.colab import files
uploaded = files.upload()
```

Choose files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving UCI\_Credit\_Card.csv to UCI\_Credit\_Card.csv

```
[ ] df = pd.read_csv(io.BytesIO(uploaded['UCI_Credit_Card.csv']))
print(df.shape)
df.head()
df
```

(30000, 25)

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6	default.payment.next.month
0	1	20000.0	2	2	1	24	2	2	-1	-1	...	0.0	0.0	0.0	0.0	689.0	0.0	0.0	0.0	0.0	1
1	2	120000.0	2	2	2	26	-1	2	0	0	...	3272.0	3455.0	3261.0	0.0	1000.0	1000.0	1000.0	0.0	2000.0	1
2	3	90000.0	2	2	2	34	0	0	0	0	...	14331.0	14948.0	15549.0	1518.0	1500.0	1000.0	1000.0	1000.0	5000.0	0
3	4	50000.0	2	2	1	37	0	0	0	0	...	28314.0	28959.0	29547.0	2000.0	2019.0	1200.0	1100.0	1069.0	1000.0	0
4	5	50000.0	1	2	1	57	-1	0	-1	0	...	20940.0	19146.0	19131.0	2000.0	36681.0	10000.0	9000.0	689.0	679.0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
29995	29996	220000.0	1	3	1	39	0	0	0	0	...	88004.0	31237.0	15980.0	8500.0	20000.0	5003.0	3047.0	5000.0	1000.0	0
29996	29997	150000.0	1	3	2	43	-1	-1	-1	-1	...	8979.0	5190.0	0.0	1837.0	3526.0	8998.0	129.0	0.0	0.0	0
29997	29998	30000.0	1	2	2	37	4	3	2	-1	...	20878.0	20582.0	19357.0	0.0	0.0	22000.0	4200.0	2000.0	3100.0	1
29998	29999	80000.0	1	3	1	41	1	-1	0	0	...	52774.0	11855.0	48944.0	85900.0	3409.0	1178.0	1926.0	52964.0	1804.0	1
29999	30000	50000.0	1	2	1	46	0	0	0	0	...	36535.0	32428.0	15313.0	2078.0	1800.0	1430.0	1000.0	1000.0	1000.0	1

30000 rows x 25 columns

```
[ ] print(df.shape)
print(df.head())
print(df.columns)
```

(30000, 25)

	ID	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	\
0	1	20000.0	2	2	1	24	2	2	-1	-1	\
1	2	120000.0	2	2	2	26	-1	2	0	0	
2	3	90000.0	2	2	2	34	0	0	0	0	
3	4	50000.0	2	2	1	37	0	0	0	0	
4	5	50000.0	1	2	1	57	-1	0	-1	0	
...	...	...	...	...	...	...	...	...	...	...	
0	...	0.0	0.0	0.0	0.0	0.0	689.0	0.0			
1	...	3272.0	3455.0	3261.0	0.0	1000.0	1000.0				
2	...	14331.0	14948.0	15549.0	1518.0	1500.0	1000.0				
3	...	28314.0	28959.0	29547.0	2000.0	2019.0	1200.0				
4	...	20940.0	19146.0	19131.0	2000.0	36681.0	10000.0				
		PAY_AMT4	PAY_AMT5	PAY_AMT6	default.payment.next.month						
0	0.0	0.0	0.0		1						
1	1000.0	0.0	2000.0		1						
2	1000.0	1000.0	5000.0		0						
3	1100.0	1069.0	1000.0		0						
4	9000.0	689.0	679.0		0						

5 rows x 25 columns

Index(['ID', 'LIMIT\_BAL', 'SEX', 'EDUCATION', 'MARRIAGE', 'AGE', 'PAY\_0', 'PAY\_2', 'PAY\_3', 'PAY\_4', 'PAY\_5', 'PAY\_6', 'BILL\_AMT1', 'BILL\_AMT2', 'BILL\_AMT3', 'BILL\_AMT4', 'BILL\_AMT5', 'BILL\_AMT6', 'PAY\_AMT1', 'PAY\_AMT2', 'PAY\_AMT3', 'PAY\_AMT4', 'PAY\_AMT5', 'PAY\_AMT6', 'default.payment.next.month'], dtype='object')

```
[ ] # Check for missing values
print(df.isnull().sum())
```

(30000, 25)

ID	0
LIMIT_BAL	0
SEX	0
EDUCATION	0
MARRIAGE	0
AGE	0
PAY_0	0
PAY_2	0
PAY_3	0
PAY_4	0
PAY_5	0
PAY_6	0
BILL_AMT1	0
BILL_AMT2	0
BILL_AMT3	0
BILL_AMT4	0
BILL_AMT5	0
BILL_AMT6	0
PAY_AMT1	0
PAY_AMT2	0
PAY_AMT3	0
PAY_AMT4	0

```
PAY_AMT5          0
PAY_AMT6          0
default.payment.next.month  0
dtype: int64
```

```
[ ] # Split the data into features (X) and target (y)
X = df.drop('default.payment.next.month', axis=1)
y = df[['default.payment.next.month']]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
[ ] sc = StandardScaler()
```

```
[ ] X_train_scaled = sc.fit_transform(X_train)
X_test_scaled = sc.fit_transform(X_test)
```

```
[ ] # Logistic Regression Model
LR = LogisticRegression()
LR.fit(X_train_scaled, y_train)
```

```
⚙ /usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1183: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for
```

```
LogisticRegression
LogisticRegression()
```

```
⚙ y_pred_lr = LR.predict(X_test_scaled)
y_pred_lr
```

```
⚙ array([0, 0, 0, ..., 0, 0, 0])
```

```
[ ] lr_cm = confusion_matrix(y_test, y_pred_lr)
lr_cm
```

```
⚙ array([[4540, 147],
       [ 994, 319]])
```


```
[ ] lr_acc_score = accuracy_score(y_test, y_pred_lr)
```

```
[ ] # Calculate precision, recall, f-score, and support
precision, recall, f_score, support = precision_recall_fscore_support(y_test, y_pred_lr, average='binary')
```

```
[ ] y_prob_lr = LR.predict_proba(X_test)[:, 1] # Probability of class 1 (default)
logloss = log_loss(y_test, y_prob_lr)
```

```
⚙ /usr/local/lib/python3.10/dist-packages/sklearn/base.py:458: UserWarning: X has feature names, but LogisticRegression was fitted without feature names
warnings.warn()
```

```
[ ] plt.figure(figsize=(8, 6))
sns.heatmap(lr_cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
[ ] print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F-score: {f_score:.4f}')
print(f'Support: {support}')
print(f'Log Loss: {logloss:.4f}')
```

```
⚙ Precision: 0.6845
Recall: 0.2430
F-score: 0.3586
Support: None
Log Loss: 7.8876
```

## ✓ KNN Classification Model

```
[ ] k = 5 # Number of neighbors
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train_scaled, y_train)
```

```
⚙ /usr/local/lib/python3.10/dist-packages/sklearn/neighbors/_classification.py:233: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
```

```
KNeighborsClassifier
```

```
KNeighborsClassifier()
```


```
[ ] y_pred_knn = knn.predict(X_test_scaled)

[ ] knn_cm = confusion_matrix(y_test, y_pred_knn)

[ ] knn_acc_score = accuracy_score(y_test, y_pred_knn)

[ ] knn_class_report = classification_report(y_test, y_pred_knn, target_names=['Non-Default', 'Default'])

[ ] plt.figure(figsize=(8, 6))
sns.heatmap(knn_cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



```
[ ] print(knn_class_report)
```

```
precision    recall  f1-score   support

Non-Default    0.84    0.92    0.88     4687
Default         0.55    0.35    0.43     1313

accuracy              0.69    0.64    0.80     6000
macro avg           0.69    0.64    0.65     6000
weighted avg        0.77    0.80    0.78     6000
```

## DecisionTree Classifier

```
[ ] dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train_scaled, y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```


```
[ ] y_pred_dt = dt.predict(X_test_scaled)


[ ] dt_cm = confusion_matrix(y_test, y_pred_dt)

[ ] dt_acc_score = accuracy_score(y_test, y_pred_dt)

[ ] dt_class_report = classification_report(y_test, y_pred_dt, target_names=['Non-Default', 'Default'])

[ ] plt.figure(figsize=(8, 6))
sns.heatmap(dt_cm, annot=True, fmt='d', cmap='Blues', cbar=False)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```





```
[ ] print(dt_class_report)
```

	precision	recall	f1-score	support
Non-Default	0.83	0.79	0.81	4687
Default	0.36	0.43	0.39	1313
accuracy			0.71	6000
macro avg	0.60	0.61	0.60	6000
weighted avg	0.73	0.71	0.72	6000

```
[ ] print("-----ACCURACY SCORE REPORT-----")
print(f'Accuracy score of Logistic Regression = {round(lr_acc_score,4)}')
print(f'Accuracy score of KNeighbors Classifier = {round(knn_acc_score,4)}')
print(f'Accuracy score of DesicionTree Classifier = {round(dt_acc_score,4)}')
```

```
-----ACCURACY SCORE REPORT-----
Accuracy score of Logistic Regression = 0.8098
Accuracy score of KNeighbors Classifier = 0.7962
Accuracy score of DesicionTree Classifier = 0.7098
```

```
[ ] Start coding or generate with AI.
```